Scalable Collaborative Filtering with Jointly Derived Neighborhood Interpolation Weights*

Robert M. Bell and Yehuda Koren

AT&T Labs – Research

180 Park Ave, Florham Park, NJ 07932

{rbell,yehuda}@research.att.com

Abstract

Recommender systems based on collaborative filtering predict user preferences for products or services by learning past user-item relationships. A predominant approach to collaborative filtering is neighborhood based ("k-nearest neighbors"), where a user-item preference rating is interpolated from ratings of similar items and/or users. We enhance the neighborhood-based approach leading to substantial improvement of prediction accuracy, without a meaningful increase in running time. First, we remove certain so-called "global effects" from the data to make the ratings more comparable, thereby improving interpolation accuracy. Second, we show how to simultaneously derive interpolation weights for all nearest neighbors, unlike previous approaches where each weight is computed separately. By globally solving a suitable optimization problem, this simultaneous interpolation accounts for the many interactions between neighbors leading to improved accuracy. Our method is very fast in practice, generating a prediction in about 0.2 milliseconds. Importantly, it does not require training many parameters or a lengthy preprocessing, making it very practical for large scale applications. Finally, we show how to apply these methods to the perceivably much slower user-oriented approach. To this end, we suggest a novel scheme for low dimensional embedding of the users. We evaluate these methods on the Netflix dataset, where they deliver significantly better results than the commercial Netflix Cinematch recommender system.

1 Introduction

Recommender systems analyze patterns of user interest in items or products to provide personalized recommendations of items that will suit a user's taste [1]. Increasingly, their excellent ability to characterize and recommend items within huge collections represent a computerized alternative to human recommendations. Because good personalized recommendations can add another dimension to the user experience, e-commerce leaders like Amazon.com and online movie rental company Netflix, have made recommender systems a salient part of their web sites.

Broadly speaking, recommender systems use either of two strategies. The *content based approach* profiles each user or product, allowing programs to associate users with matching products. For example, a movie profile might describe its genre, the participating actors, its box office popularity, etc. User profiles could include demographic information or answers to a suitable questionnaire. Of course, content based strategies require gathering external information that might not be available or easy to collect.

We focus on an alternative strategy, known as *Collaborative Filtering* (CF) [6], which relies only on past user behavior—e.g., their previous transactions or product ratings—and does not require the creation of explicit profiles. CF analyzes relationships between users and interdependencies among products, in order to identify new user-item associations. Besides avoiding the need for extensive data collection about items or users, CF requires no domain knowledge. In addition, it offers the potential to uncover patterns that would be difficult or impossible to profile using content based techniques. This has led to many papers (e.g., [8]), research projects (e.g., [9]) and commercial systems (e.g., [11]) based on CF.

From a more abstract perspective, CF can be cast as missing value estimation. Given a user-item matrix of scores with many missing values, our goal is to estimate the missing values based on the given ones. The observed user-item scores measure the amount of interest between respective users and items. We call these user-item scores *ratings*, and they constitute the input to our algorithm.

The most common form of CF is the neighborhood-based approach (also known as "k Nearest Neighbors" or kNN, for short). The kNN methods identify pairs of items that tend to be rated similarly or like-minded users with similar histories of rating or purchasing, in order to predict ratings for unobserved users-item pairs.

Three major components characterize the kNN approach: (1) data normalization, (2) neighbor selection, and

---

(3) determination of interpolation weights. Our experience has shown little difference among neighbor selection strategies (e.g., distance-based vs. correlation-based). However, the two other components, namely data normalization and interpolation weights, have proved vital to the success of the scheme. Accordingly, we revisit these two components and suggest novel methods to significantly improve the accuracy of kNN approaches without meaningfully affecting running time. Our three main contributions are:

1. It is customary to normalize the data before activating a kNN method. This brings different ratings to a closer level, which allows a better mixing thereof. Usually this is achieved by adjusting for the varying mean ratings across users and/or items. In Section 3 we offer a more comprehensive normalization that considers additional effects that are readily available in virtually all ratings datasets. This allows us to explain and eliminate much of the interfering variability from the data and to work with residuals that are more amenable to mutual interpolation.

2. Past kNN methods relate items (or users) by various heuristic variants of correlation coefficients, which allowed direct interpolation from neighbors' scores. In Section 4 we offer a rigorous alternative to these interpolation weights based on global optimization of a cost function pertaining to all weights simultaneously. This results in another improvement of estimation quality with a minor increase in running time.

3. The kNN approach can capitalize on two different kinds of information: item-item, how the subject user rated similar items; and user-user, how the subject item was rated by like-minded users. Because there typically are many more users than items in a system, the user-oriented approach is known as slower and less accurate. In Section 5 we discuss benefits of also employing a user-oriented approach, and we offer a novel method for alleviating the inherent computational difficulties of the user-oriented approach.

An encouraging evaluation of the results on the Netflix Prize user-movie dataset [3] is provided in Section 6.

2 Related Works

We are given ratings about m users and n items, arranged in an $m \times n$ matrix $R = \{r_{ui}\}_{1 \leqslant u \leqslant m, 1 \leqslant i \leqslant n}$. We reserve special indexing letters for distinguishing users from items: for users u, v, and for items i, j, k.

2.1 Neighborhood-based collaborative filtering

The original neighborhood-based approach, which was shared by virtually all earlier CF systems, is user-oriented;

see [8] for a good analysis. In order to estimate the unknown rating r_{ui}, we resort to a set of *users* $\mathrm{N}(u; i)$ that tend to rate similarly to u ("neighbors"), and that actually rated item i (i.e., r_{vi} is known for each $v \in \mathrm{N}(u; i)$).

The predicted value of r_{ui} is taken as a weighted average of the neighbors' ratings:

$$r_{ui} \leftarrow \frac{\sum_{v \in \mathrm{N}(u;i)} s_{uv} r_{vi}}{\sum_{v \in \mathrm{N}(u;i)} s_{uv}} \qquad (1)$$

An analogous alternative to the user-oriented approach is the item-oriented approach [11, 15]. In those methods, to estimate the unknown r_{ui}, we identify a set of neighboring *items* $\mathrm{N}(i; u)$ that other users tend to rate similarly to their rating of i. All items in $\mathrm{N}(i; u)$ must have been rated by u. Then, in parallel to (1), the estimated value of r_{ui} is taken as a weighted average of the ratings of neighboring items:

$$r_{ui} \leftarrow \frac{\sum_{j \in \mathrm{N}(i;u)} s_{ij} r_{uj}}{\sum_{j \in \mathrm{N}(i;u)} s_{ij}} \qquad (2)$$

The similarities—denoted by s_{uv} or s_{ij}—play a central role here, as they are used both for selecting the neighbors and for weighting the above averages. Common choices are the Pearson correlation coefficient and the closely related cosine similarity. Methods also differ by how they normalize or center data prior to activating interpolation rule (1) or (2). For example, correcting for user-specific means can improve prediction quality. We present a more comprehensive treatment of data normalization in Section 3. Sarwar et al. [15] found that item-oriented approaches deliver better quality estimates than user-oriented approaches while allowing more efficient computations. The greater efficiency occurs because, typically, the number of items is significantly lower than the number of users, which allows precomputing all item-item similarities for retrieval as needed.

Neighborhood-based methods became very popular because they are intuitive and relatively simple to implement. In particular, they do not require tuning many parameters or an extensive training stage. They also provide a concise and intuitive justification for the computed predictions. This enables presenting the user a list of similar items that he or she has previously rated, as the basis for the estimated rating. This way, the user actually understands the reasoning behind the recommendation, allowing him/her to better assess its relevance (e.g., downgrade the estimated rating if it is based on an item that he or she no longer likes), or even encourage the user to alter outdated ratings.

However, standard neighborhood-based methods raise some concerns:

1. The similarity function (s_{uv} or s_{ij}), which directly defines the interpolation weights, is arbitrary. Various CF algorithms use somewhat different similarity measures, trying to quantify the elusive notion of user- or

item-similarity. Suppose that a particular item is predicted perfectly by a subset of the neighbors. In that case, we would want the predictive subset to receive all the weight, but that is impossible for bounded similarity scores like the Pearson correlation coefficient.

2. Previous neighborhood-based methods do not account for interactions among neighbors. Each similarity between an item i and a neighbor $j \in N(i; u)$ is computed independently of the content of $N(i; u)$ and the other similarities: s_{ik} for $k \in N(i; u) - \{j\}$. For example, suppose that our items are movies, and the neighbors set contains three movies that are highly correlated with each other (e.g., sequels such as "Lord of the Rings 1–3"). An algorithm that ignores the similarity of the three movies when determining their interpolation weights, may end up essentially triple counting the information provided by the group.

3. By definition, the interpolation weights sum to one, which may cause overfitting. Suppose that an item has no useful neighbors rated by a particular user. In that case, it would be best to ignore the neighborhood information, staying with the current data normalization. Nevertheless, the standard neighborhood formula uses a weighted average of ratings for the uninformative neighbors.

4. Neighborhood methods may not work well if variability differs substantially among neighbors.

In a recent work [2], we overcame most of these shortcomings, but the resulting algorithm was several orders of magnitude slower than previous neighborhood-methods, thereby limiting its practical applicability. In this work, we address the aforementioned issues of neighborhood based approaches, without compromising running time efficiency.

3 Normalizing by Removing Global Effects

Although neighborhood and factorization based CF are both powerful prediction methods, there are several reasons to precede either technique with simple models that estimate what we call "global effects." First, there may be large user and item effects—i.e., systematic tendencies for some users to give higher ratings than other users and for some items to receive higher ratings than others. The basic kNN interpolation method detailed in equations (1)-(2) requires ratings where user and item effects have been taken out in order to, e.g., avoid predicting too high for low-rated items that happen to have a lot of neighbors with high average ratings, and vice versa.

Second, one may have access to information about either the items or users that can benefit the model. Although factorization offers the potential to detect such structure through estimation of latent variables, directly incorporating variables such as movie genre and user demographics

may be more effective and does not require training a factorization model. While such content based analysis is beyond the scope of this paper, we do consider two types of easily derived variables: the number of ratings of an item or by a user and the average rating of an item or by a user. Variables like these allow us, for example, to distinguish users who like the most commonly rated movies best from those who prefer more specialized fare (after controlling for both movie and user effects).

Third, there may be characteristics of specific ratings, such as the date of the rating, that explain some of the variation in scores. For example, a particular user's ratings may slowly, or suddenly, rise over time, above and beyond any change explained by the inherent quality of the items being rating. Similarly, ratings for some movies may fall with time after their initial release dates, while others stand the test of time quite well. Neither factorization nor kNN could be expected to detect patterns like these.

3.1 Methods

Our strategy is to estimate one "effect" at a time, in sequence (i.e., the main effect for items, the main effect for users, a user-time interaction, etc.). At each step, we use residuals from the previous step as the dependent variable for the current step. Consequently, after the first step, the r_{ui} refer to residuals, rather than raw ratings.

For each of the effects mentioned above, our goal is to estimate either one parameter for each item or one parameter for each user. For the rest of this subsection, we describe our methods for estimating user specific parameters; the method for items is perfectly analogous.

We denote by x_{ui} the explanatory variable of interest corresponding to user u and item i. For user main effects, the x_{ui}'s are identically 1. For other global effects, we center x_{ui} for each user by subtracting the mean of x_{ui} for that user. In each case, our model is:

$$r_{ui} = \theta_u x_{ui} + error \qquad (3)$$

With sufficient ratings for user u, we might use the unbiased estimator:

$$\hat{\theta}_u = \frac{\sum_i r_{ui} x_{ui}}{\sum_i x_{ui}{}^2}$$

where each summation is over all items rated by u. However, for sparse data, some values of $\hat{\theta}_u$ may be based on very few observations, resulting in unreliable estimates.

To avoid overfitting, we shrink individual values of $\hat{\theta}_u$ towards a common value. Shrinkage can be motivated from a Bayesian perspective. Suppose that the true θ_u are independent random variables drawn from a normal distribution,

$$\theta_u \sim N(\mu, \tau^2)$$

for known μ and τ^2, while

$$\hat{\theta}_u | \theta_u \sim N(\theta_u, \sigma_u^2)$$

for known σ_u^2. We estimate θ_u by its posterior mean:

$$E(\theta_u | \hat{\theta}_u) = \frac{\tau^2 \hat{\theta}_u + \sigma_u^2 \mu}{\tau^2 + \sigma_u^2}$$

a linear combination of the empirical estimator $\hat{\theta}_u$ and the common mean μ. The parameter σ_u^2 can be estimated from the formula for the variance of a weighted mean, while μ can be estimated by the mean of the θ_u's (optionally weighted by n_u). Empirical Bayes [4] suggests that the maximum likelihood estimate of τ^2 can be found as the solution to:

$$\tau^2 = \frac{\sum_u [(\hat{\theta}_u - \mu)^2 - \sigma_u^2]/(\tau^2 + \sigma_u^2)^2}{\sum_u (\tau^2 + \sigma_u^2)^{-2}}$$

In practice, we used a slightly simpler estimator for θ_u by assuming $\mu = 0$ and σ_u^2 is proportional to $1/n_u$, to yield:

$$\frac{n_u \hat{\theta}_u}{n_u + \alpha}$$

where n_u is the number of ratings by user u and α is a constant. We determined α by cross validation.

3.2 Example

We illustrate the impact of estimating successive global effects on the Netflix data [3]. The dataset is based on more than 100 million ratings of movies performed by anonymous Netflix customers and is described in details in Section 6. Here, we report results on the *Probe set*, which is a test set containing about 1.4 million ratings compiled by Netflix. Accuracy of predictions is measured by their root mean squared error (RMSE).

Table 1 shows how the RMSE for the probe set declines with the inclusion of each successive global effect. Not surprisingly, by far the largest improvements in RMSE are associated with the two sets of main effects for movies and for users. They reduce the RMSE from 1.1296 based on use of the mean rating in the training data, down to 0.9841.

Of more interest are various interactions. The first interaction term, User×Time(user)$^{1/2}$, allows each user's rating to change linearly with the square root of the number of days since the user's first rating. Exploratory analysis indicated that this transformation improved the fit relative to the untransformed number of days. Technically, we set each x_{ui} in (3) to the square root of the number of days elapsed since the first rating by user u till the time u rated item i. Then, for each fixed user u, we center all computed values of x_{ui}, and calculate $\hat{\theta}_u$ in order to predict r_{ui} based on x_{ui}. This interaction, reduces the RMSE by 0.0032.

Similarly, User×Time(movie)$^{1/2}$, allows each user's rating to change linearly with the square root of the number of days since the movie's first rating by anyone. Somewhat surprisingly, this latter interaction contributes almost

Effect	RMSE	Improvement
Overall mean	1.1296	NA
Movie effect	1.0527	.0769
User effect	0.9841	.0686
User×Time(user)$^{1/2}$	0.9809	.0032
User×Time(movie)$^{1/2}$	0.9786	.0023
Movie×Time(movie)$^{1/2}$	0.9767	.0019
Movie×Time(user)$^{1/2}$	0.9759	.0008
User×Movie average	0.9719	.0040
User×Movie support	0.9690	.0029
Movie×User average	0.9670	.0020
Movie×User support	0.9657	.0013

Table 1. RMSE for Netflix probe data after adding a series of global effects to the model

as much as the first. Two additional time interactions concentrate on the complementary movie viewpoint. That is, for each movie, they model how its ratings change with time. Once again, the time variables are either square root of days since first rating of the movie, or square root of days since first rating by the user.

The next two effects interact users with movie characteristics. We measure the popularity of a movie in two different ways: (1) its average rating; and (2) its support, which is the number of ratings associated with the movie. These effects—User×Movie average and User×Movie support—measure how users change their ratings based on the popularity of the movies. We try to estimate here how closely a user follows the public opinion, or perhaps the opposite, how contrarian is the user. The most effective interaction is between users and movie average, indicating that users varied in terms of how much they agree with the consensus.

Finally, we interact movies with user characteristics, regressing the ratings of each movie on the mean or support of the users. Although these interactions are more difficult to motivate, each contributes an additional modest decrease of the RMSE. Overall, the eight interactions combine to reduce the RMSE by 0.0184 to 0.9657.

4 A Neighborhood Relationships Model

In this section we assume that global effects have already been removed, so whenever we refer to a rating we actually mean the residual remaining after removing global effects from the original ratings data. However, our discussion here is completely general, so the following method applies whether global effects are removed or not.

We need to estimate the unknown rating by user u of item i, that is r_{ui}. As with all neighborhood-based methods, our first step is neighbor selection. In this section, we focus on an item-oriented method. Among all items rated by u, we select the K most similar to i, N$(i; u)$, by using a similarity function such as the correlation coefficient, as described later in this section. Typical values of K lie in the

range of 20–50; see Section 6.

Given a set of neighbors $N(i; u)$, we need to compute *interpolation weights* $\{w_{ij} | j \in N(i; u)\}$ that will enable the best prediction rule of the form:

$$r_{ui} \leftarrow \sum_{j \in N(i;u)} w_{ij} r_{uj} \qquad (4)$$

For notational convenience assume that the K neighbors in $N(i; u)$ are indexed by $1, \ldots, K$.

We seek a formal computation of the interpolation weights, that stems directly from their usage within prediction rule (4). As explained earlier, it is important to derive all interpolation weights simultaneously to account for interdependencies among the neighbors. We achieve these goals by defining a suitable optimization problem.

To start, we consider a hypothetical dense case, where all users but u rated both i and *all* its neighbors in $N(i; u)$. In that case, we could learn the interpolation weights by modeling the relationships between item i and its neighbors through a least squares problem:

$$\min_w \sum_{v \neq u} \left(r_{vi} - \sum_{j \in N(i;u)} w_{ij} r_{vj} \right)^2 \qquad (5)$$

Notice that the only unknowns here are the w_{ij}'s. The optimal solution to the least squares problem (5) is found by differentiation as a solution of a linear system of equations. From a statistics viewpoint, it is equivalent to the result of a linear regression (without intercept) of r_{vi} on the r_{vj} for $j \in N(i; u)$. Specifically, the optimal weights are given by:

$$Aw = b \qquad (6)$$

Here, A is a $K \times K$ matrix defined as:

$$A_{jk} = \sum_{v \neq u} r_{vj} r_{vk} \qquad (7)$$

Similarly the vector $b \in \mathbb{R}^K$ is given by:

$$b_j = \sum_{v \neq u} r_{vj} r_{vi} \qquad (8)$$

For a sparse ratings matrix there are likely to be very few users who rated i and all its neighbors $N(i; u)$. Accordingly, it would be unwise to base A and b as given in (7)–(8) only on users with complete data. Even if there are enough users with complete data for A to be nonsingular, that estimate would ignore a large proportion of the information about pairwise relationships among ratings by the same user. However, we can still estimate A and b, up to the same constant, by averaging over the given pairwise support, leading to the following reformulation:

$$\bar{A}_{jk} = \frac{\sum_{v \in U(j,k)} r_{vj} r_{vk}}{|U(j,k)|} \qquad (9)$$

$$\bar{b}_j = \frac{\sum_{v \in U(i,j)} r_{vj} r_{vi}}{|U(i,j)|} \qquad (10)$$

where $U(j, k)$ is the set of users who rated both j and k.

This is still not enough to overcome the sparseness issue. The elements of $\bar{A}_{jk}$ or $\bar{b}_j$ may differ by orders of magnitude in terms of the number of users included in the average. As discussed in the previous section, averages based on relatively low support (small values of $|U(j, k)|$) can generally be improved by shrinkage towards a common mean. Thus, we compute a baseline value that is defined by taking the average of all possible $\bar{A}_{jk}$ values. Let us denote this baseline value by avg; its precise computation is described in the next subsection. Accordingly, we define the corresponding $K \times K$ matrix $\hat{A}$ and the vector $\hat{b} \in \mathbb{R}^K$:

$$\hat{A}_{jk} = \frac{|U(j,k)| \cdot \bar{A}_{jk} + \beta \cdot avg}{|U(j,k)| + \beta} \qquad (11)$$

$$\hat{b}_j = \frac{|U(i,j)| \cdot \bar{b}_j + \beta \cdot avg}{|U(i,j)| + \beta} \qquad (12)$$

The parameter β controls the extent of the shrinkage. A typical value on the Netflix data when working with residuals of full global effects is $\beta = 500$.

Our best estimate for A and b are $\hat{A}$ and $\hat{b}$, respectively. Therefore, we modify (6) so that the interpolation weights are defined as the solution of the linear system:

$$\hat{A}w = \hat{b} \qquad (13)$$

The resulting interpolation weights are used within (4) in order to predict r_{ui}. When working with residuals of global effects, they should be added back to the predicted r_{ui}, which completes the computation.

This method addresses all four concerns raised in Section 2.1. First, interpolation weights are derived directly from the ratings, not based on any similarity measure. Second, the interpolation weights formula explicitly accounts for relationships among the neighbors. Third, the sum of the weights is not constrained to equal one. If an item (or user) has only weak neighbors, the estimated weights may all be very small. Fourth, the method automatically adjusts for variations among items in their means or variances.

4.1 Preprocessing

Efficient computation of an item-item neighborhood-based method requires precomputing certain values associated with each item-item pair for rapid retrieval.

First, we need a quick access to all item-item similarities, the s_{ij} values. These similarities are required for identifying the K neighbors that constitute $N(i; u)$. Here, we usually take s_{ij} as the Pearson correlation coefficient between

i and j calculated on their common raters. It is important to shrink s_{ij} based on their support, e.g., multiplying the correlation by: $|U(i,j)|/(|U(i,j)| + \alpha)$ for some small α.

Second, we precompute all possible entries of $\hat{A}$ and $\hat{b}$. To this end, for each two items i and j, we compute:

$$\bar{A}_{ij} = \frac{\sum_{v \in U(i,j)} r_{vi} r_{vj}}{|U(i,j)|}$$

Then, the aforementioned baseline value avg, which is used in (11)-(12), is taken as the average entry of the precomputed $n \times n$ matrix $\bar{A}$. In fact, we recommend using two different baseline values, one by averaging the non-diagonal entries of $\bar{A}$ and another one by averaging the generally-larger diagonal entries. Finally, we derive a full $n \times n$ matrix $\hat{A}$ from $\bar{A}$ by (11), using the appropriate value of avg. Here, the non-diagonal average is used when deriving the non-diagonal entries of $\hat{A}$, whereas the diagonal average is used when deriving the diagonal entries of $\hat{A}$.

Because of symmetry, it is sufficient to store the values of s_{ij} and $\hat{A}_{ij}$ only for $i \geqslant j$. We allocated one byte for each individual value, so the overall space required for n items is exactly $n(n + 1)$ bytes. For example, for the Net-flix dataset which contains 17770 movies, overall memory requirements are 300MB, easily fitting within core memory. A more comprehensive system that includes data on 100,000 items would require about 10GB of space, which can fit within core memory of current 64bit servers. For even larger systems, disk resident storage of item-item values is still practical, as evident by the Amazon item-item recommender system, which accesses stored similarity information for several million catalog items [11]. Preprocessing time is quadratic in n and linear in the number of ratings. The time required for computing all s_{ij}'s and $\hat{A}_{ij}$'s on the Netflix data (which contains 100 million ratings) was about 15 minutes on a Pentium 4 PC. Notice that preprocessing can be easily parallelized.

Precomputing all possible entries of matrix $\hat{A}$ saves the otherwise lengthy time needed to construct entries on the fly. After quickly retrieving the relevant entries of $\hat{A}$, we can compute the interpolation weights by solving a $K \times K$ system of equations. For typical values of K (between 20 and 50), this time is comparable to the time needed for computing the K nearest neighbors, which is common to all neighborhood-based approaches. Hence, while our method relies on a much more detailed computation of the interpolation weights compared to previous methods, it does not significantly increase running time; see Section 6.

4.2 Calculation of the weights

The interpolation weights can be computed by solving (13) using standard linear equations solvers. However, we experience a modest increase in prediction accuracy when w is constrained to be nonnegative. This requires a

---

NonNegativeQuadraticOpt $(A \in \mathbb{R}^{K \times K}, b \in \mathbb{R}^K)$
% Minimize $x^T A x - 2b^T x$ s.t. $x \geqslant 0$

> **do**
> $r \leftarrow b - Ax$ *% the residual, or "steepest gradient"*
> *% find active variables - those that are pinned due to*
> *% nonnegativity constraints; set respective r_i's to zero*
> **for** $i = 1, \ldots, k$ **do**
>   **if** $x_i = 0$ and $r_i < 0$ **then**
> $r_i \leftarrow 0$
>   **end if**
> **end for**
> $\alpha \leftarrow \frac{r^T r}{r^T A r}$ *% max step size*
> *% adjust step size to prevent negative values:*
> **for** $i = 1, \ldots, k$ **do**
>   **if** $r_i < 0$ **then**
> $\alpha \leftarrow \min(\alpha, -x_i/r_i)$
>   **end if**
> **end for**
> $x \leftarrow x + \alpha r$
> **while** $\|r\| < \epsilon$ *% stop when residual is close to 0*
> **return** x

Figure 1. Minimizing a quadratic function with non-negativity constraints

quadratic program, which can be solved by calling "Non-NegativeQuadraticOpt($\hat{A}$, $\hat{b}$)", as shown in Figure 1. The function is based on the principles of the Gradient Projection method; see, e.g., [12].

5 User-oriented computation

A user-oriented approach can be implemented by switching the roles of users and items throughout the derivation of the previous section. However, as with previous neighborhood-based approaches, an item-oriented approach enables a much faster implementation by precomputing and storing in main memory a full item-item matrix containing all s_{ij} and $\hat{A}_{ij}$ values for future retrieval. Typically, the larger number of users complicates such a precomputation in terms of both time and space, and out-of-core storage is required. Moreover, our experience with the Netflix data, under various settings, showed that an item-oriented approach consistently delivered more accurate predictions than a user-oriented one.

Nonetheless, user-oriented approaches identify different kinds of relations that item-oriented approaches may fail to recognize, and thus can be useful on certain occasions. For example, suppose that we want to predict r_{ui}, but none of the items rated by user u is really similar to i. In this case, an item-oriented approach will face obvious difficulties. However, when employing a user-oriented perspective, we may find a set of users similar to u, who rated i. The ratings of i

by these users would allow us to improve prediction of r_{ui}.

Another common case is when a user u has barely provided any ratings to the system, but we have a rich history of past transactions by him—e.g., his purchase history, viewed pages history, searched items and other kinds of implicit information. By employing a user-oriented method on the implicit information, we can relate user u to other users who did provide ratings for item i, thereby derive r_{ui}. Finally, we can improve prediciton accuracy by mixing the results of the item-oriented approach with those of the user-oriented one. In this section, we present a way to efficiently capitalize on the additional information captured within user-user relationships.

5.1 Low dimensional embedding of the users

When applying our method in a user-oriented manner, the major computational bottleneck is the creation of N$(u; i)$, the K users most similar to u that rated item i. Here, the typically huge number of user-user similarities complicates their storage and precomputation. Moreover, because we expect many more users than items, scanning all users that rated i for picking the K most similar to u becomes a rather lengthy process. Hence, it is crucial to measure the similarity between users very efficiently.

We significantly lower the computational effort of measuring user-user similarities by embedding all users in a low dimensional Euclidean space. To this end, we view each user as a point within the space of n movies. A common way for reducing dimensionality is by using Principal Component Analysis (PCA), which is often carried out through the closely related Singular Value Decomposition (SVD). Recall that the n-dimensional user points are arranged within the $m \times n$ matrix R. Consequently, SVD computes the best rank-f approximation R^f, which is defined as the product of two rank-f matrices $P_{m \times f}$ and $Q_{n \times f}$, where $\|R - PQ^T\|_F$ is minimized. In a way, the matrix P is an optimal f-dimensional embedding of the users.

Applying an SVD-based technique to CF raises unique difficulties due to the sparsity issue. The conventional SVD computation requires that all entries of R are known. In fact, the goal of SVD is not properly defined when some entries of R are missing. Previous works have adapted matrix factorizations techniques to handle missing data in a variety of ways [2, 5, 7, 13, 14]. For our special needs, we suggest the following approach.

Let us symbolize the set of (u, i) pairs for which r_{ui} is known by $\mathcal{K}$. The goal of SVD, when restricted to the known ratings, is to minimize:

$$\sum_{(u,i) \in \mathcal{K}} (r_{ui} - p_u^T q_i)^2 \qquad (14)$$

Here, p_u is the u-th row of P, which corresponds to user u. Likewise, q_i is the i-th row of Q, which corresponds to

item i. A critical issue is to avoid overfitting for items and users with relatively sparse data. To this end we regularize the model by penalizing the norm of each p_u and q_i ("ridge regression"). That is, we replace (14) with:

$$\sum_{(u,i) \in \mathcal{K}} (r_{ui} - p_u^T q_i)^2 + \lambda(\|p_u\|^2 + \|q_i\|^2) \qquad (15)$$

A typical choice of λ for the Netflix data is 0.05. To minimize (15), we employ an alternating least squares strategy. That is, we alternate between fixing Q and P, thereby obtaining a series of efficiently solvable least squares problems. To further avoid overfitting, we restrict all entries of P and Q to be nonnegative by using a *nonnegative least squares* solver [10]. In essence, P and Q form a regularized nonnegative matrix factorization for the partially observed matrix R.

When recomputing P, we address each user u as follows. Recall that n_u is the number of ratings by u. We use the $n_u \times f$ matrix $Q[u]$ to denote the restriction of Q to the items rated by u, and let the vector $r_u \in \mathbb{R}^{n_u}$ contain the given ratings by u ordered as in $Q[u]$. The new value for p_u is given by solving a nonnegative least-squares problem:

$$\begin{pmatrix} Q[u] \\ \Lambda \end{pmatrix} p_u = \begin{pmatrix} r_u \\ 0 \end{pmatrix}$$

Here, Λ is an $f \times f$ diagonal matrix, where all diagonal entries are $\sqrt{\lambda n_u}$. The computation of Q is analogous.

Finally, following a few tens of iterations of recomputing P and Q, we converge at an f-dimensional embedding of the users - P. We use $f = 10$. Of course, the low dimensional embedding of the users is performed only once, at the preprocessing stage. Since then, whenever we need to find neighboring users, we do so in their low dimensional representation. This significantly alleviates the computational complexity of the user-based computations and facilitates their use in real life, large datasets, such as the Netflix data. Further performance gains can be achieved by organizing the f-dimensional user points within a space-partitioning data structure (such as a kd-tree), which allows very efficient retrieval of the nearest points (most similar users).

After identifying the K most similar users, N$(u; i)$, we continue with the usual computation of interpolation weights. Unlike the movie-oriented case, the sheer number of users prevents us from precomputing their inner products. However, at this stage we need to deal with only K users, which is not a major bottleneck in the process, especially when considering that individual users are typically associated with far fewer ratings compared to individual items. We provide experimental results in Subsection 6.3.

5.2 Unifying user- and item-relations in a single model

Our current model assumes that relationships among users are fixed across all items. However, in reality, a user v

may be very predictive of user u for certain kinds of items, but far less predictive for other items. When predicting r_{ui}, we would like to derive user-user interpolation weights that reflect how the neighboring users relate to u with respect to the given item – i. Thus, when learning the interpolation weights, we give a higher weight to items similar to i, which may serve as a better proxy for the sought user relations. To this end, we introduce item-item similarities (s_{ij}) into the user version of (5), which becomes:

$$\min_w \sum_{j \neq i} s_{ij} \left(r_{uj} - \sum_{v \in N(u;i)} w_{uv} r_{vj} \right)^2 \qquad (16)$$

Similarly, the user-version of matrix A and vector b of (7)-(8), becomes: $A_{v_1 v_2} = \sum_{j \neq i} s_{ij} r_{v_1 j} r_{v_2 j}$, $b_v = \sum_{j \neq i} s_{ij} r_{vj} r_{uj}$. Essentially, these modifications inject item-item relationships into the user-user model. Possible choices for s_{ij} are the absolute value of the Pearson correlation coefficient, or an inverse of the squared error. As usual with item-item magnitudes, all s_{ij}'s can be precomputed and stored, so introducing them into the user-user model barely affects running time while benefiting prediction accuracy. Naturally, this is part of our default setting.

A parallel idea can be used for integrating user-awareness into the item-oriented model. However, it requires the item-item inner products to be computed specifically for each query in order to reflect the relevant user similarities. This prevents the precomputation and storage of all item-item inner products. Because items are typically associated with large numbers of ratings, an online computation of their inner products is very expensive and impractical for large datasets. Hence, currently, we are not adding user-awareness into the item-oriented models.

6 Experimental Study

We evaluated our algorithms on the Netflix data of more than 100 million movie ratings [3]. We are not aware of any publicly available CF dataset that is close to the scope and quality of this dataset.

To maintain compatibility with results published by others, we adopted some standards that were set by Netflix, as follows. First, quality of the results is measured by their root mean squared error (RMSE), a measure that puts more emphasis on large errors compared with the alternative of mean absolute error. In addition, we report results on two test sets compiled by Netflix, one is known as *the Probe set* and the other is known as *the Quiz set*. These two test sets were constructed similarly, with both containing about 1.4 million of the most recent movie ratings (by date) performed by the users. The true ratings of the Probe set are provided. However, we do not know the true ratings for the Quiz set, which are held by Netflix being the subject of the Netflix Prize contest. Netflix provides RMSE values to

competitors that submit their predictions for the Quiz set. The benchmark is Netflix's proprietary CF system, Cinematch, which achieved a RMSE of 0.9514 on this Quiz set. The two test sets contain many more ratings by users that do not rate much and are harder to predict. In a way, they represent real requirements from a CF system, which needs to predict new ratings from older ones, and to equally address all users, not only the heavy raters.

6.1 Experiments with the Probe set

Table 2 compares our item-oriented method to a kNN method that takes interpolation weights as Pearson correlation coefficients (shrunk to improve results), which represent the common approach to the problem. Both methods use the same item neighborhoods. Reported results are RMSE values computed on the Probe set. We apply the methods after varying stages of preprocessing the data. First, we applied the methods to the raw data, where scores were not normalized at all. Then, we removed the first two effects: the user effect and the item effect, which is similar to double-centering. When no neighborhood-interpolation is performed, such a normalization alone delivers a RMSE of 0.9841, as shown in the second column of the table. A more substantial normalization is achieved by accounting for all the global effects, as described in Section 3. Global effects on their own lower the Probe RMSE to 0.9657, before applying any neighborhood interpolation.

Finally, our method can be combined with a factorization approach, which models the ratings matrix through a low rank factorization [2, 5, 7, 13, 14]. In a way, factorization complements the more localized neighborhood-based methods by providing a higher level perspective to the ratings data. Notice that unlike the previously mentioned normalizations, factorization requires an extensive training in order to learn the factors. Our implementation delivered a RMSE of 0.9167 on the Probe set, before considering any neighborhood information. A kNN method can be used to improve factorization results, by applying it to the residuals remaining after subtracting the estimates generated by factorization. Effectively, error is smoothed by considering local information that the rather regional factorization approaches tend to miss.

The remaining columns of Table 2 contain results using the two methods for determining interpolation weights. We provide results representing the spectrum of viable choices for neighborhood size – K. For each value of K, the same neighborhoods are used for the two methods; the only difference lies in the determination of interpolation weights. The main findings are as follows.

- The jointly derived interpolation weights proposed in this paper uniformly outperformed standard correlation based weights. The improvement was 0.0071 when applied after factorization, but was much larger

Data normalization	No interpolation ($k = 0$)	Correlation-based interpolation			Jointly derived interpolation		
		$k = 20$	$k = 35$	$k = 50$	$k = 20$	$k = 35$	$k = 50$
none (raw scores)	NA	0.9947	1.002	1.0085	0.9536	0.9596	0.9644
double centering	0.9841	0.9431	0.9470	0.9502	0.9216	0.9198	0.9197
global effects	0.9657	0.9364	0.9390	0.9413	0.9194	0.9179	0.9174
factorization	0.9167	0.9156	0.9142	0.9142	0.9071	0.9071	0.9071

Table 2. Comparing our interpolation scheme against conventional correlation-based interpolation, by reporting RMSEs on the Probe set. Various levels of data normalization (preprocessing) are shown, and different sizes of item-neighborhoods (K) are considered.

with less extensive data normalizations.

- Use of the neighborhood model with our proposed interpolation weights substantially improved the predictions across all the data normalizations. The neighborhood approach reduced RMSE by 0.0096 even after the much more time consuming factorization method. Notably, the correlation-based neighborhood model produced a much more modest improvement of 0.0025 relative to factorization.

- The best neighborhood size K varied depending on the extent of prior data normalization. Interestingly, we found no difference in RMSE over the range 20 to 50 neighbors for the best combination—jointly derived interpolation weights after factorization.

- Even with incorporation of neighborhood information, more complete data normalization improved predictions, but with diminishing returns. For example, while global effects reduced RMSE by 0.0184 before the neighborhood model, only 0.0023 survived after incorporation of the neighborhood model.

The correlation based method required about 200 seconds to process the whole Probe set regardless of the value of K. Our method, with $K = 20$, increased the time slightly to about 235 seconds for the whole Probe set, where the added time represents the effort needed to solve the $K \times K$ least squares problem for each query. Not surprisingly, increasing K further raised running times to around 355 seconds ($K = 35$) or 470 seconds ($K = 50$), due to the growing effort needed to solve the larger least squares problems. Even so, the slower mode of our method ($K = 50$) processed a single rating query in less than 0.4 millisecond.

6.2   Experiments with the Quiz set

When computing predictions for the Quiz set, the test set held by Netflix, we subsume the Probe set into our training data. This makes Quiz results better than the Probe results. Accordingly, our item-oriented method produced a RMSE of 0.9086 (4.5% improvement over Netflix system) for the quiz set when applied to residuals of global effects. When applied to residuals of factorization the RMSE dropped to 0.8982 (5.6% improvement over Netflix system), at the cost of requiring the training of a factorization model.

These results are comparable to recently reported results for the Netflix data by methods that require an extensive training phase [2, 13]. Moreover, when mixed with other methods, results can be substantially improved. For example, an alternative method to introducing neighborhood-awareness into factorization-based results was described in [2]. While the accuracy of the two methods is similar, they differ considerably in their nature, and thus produce different ratings. Thus, they can be mixed effectively to produce a more accurate solution. Their mixtures achieve solutions with RMSEs around 0.8900 (6.45% improvement over Netflix system), which would currently rank high on the Netflix Prize Leaderboard[1], even before mixing with the user-oriented approach to which we now turn.

6.3   Experiments with the user-oriented approach

We evaluated the user-oriented approach on the Quiz set. In contrast to the item-oriented method, which requires around 30–50 neighbors to maximize prediction accuracy, the user-oriented approach benefits from a larger number of neighbors. Our default setting uses 100 neighbors. This probably reflects the reality that the data contains many more users than movies, but individual users provide a relatively small amount of information. The low dimensional embedding of the users was essential for dealing with the Netflix data, by allowing prediction of a single rating in around 40 milliseconds. Though this is much slower than the movie-oriented computation, it is still considerably faster than the common full-dimensional computation, which requires about 400 milliseconds per prediction. Moreover, as hinted in Subsection 5.1, a more efficient utilization of the low dimensional embedding would involve a space-partitioning data structure in order to further speed up the neighborhood retrieval. Computation of the 10-dimesnisonal embedding itself required about 40 minutes.

While a low dimensional embedding of users allowed a 10X speedup compared with the full high-dimensional computation, it did not reduce prediction accuracy. The RMSE when employing the method on normalized data (after global effects were removed) is 0.9157. Surprisingly,

---

[1]www.netflixprize.com/leaderboard

this compares favorably, with a RMSE of 0.9174 achieved through a much longer run that computed neighborhoods on the original data. Hence, we deduce that dimensionality reduction is an effective tool for performing user-oriented methods.

While a RMSE of 0.9157 presents a 3.75% improvement over Netflix Cinematch result, it is still worse than our reported 0.9086 RMSE using the item-oriented method. However, since the two methods capture different sorts of information, we can get a better result by combining them. For example by taking a weighted average of the results (60% movie, 40% user), the RMSE was lowered to 0.9030. More sophisticated combinations might go beyond using a fixed ratio, to modify the ratio across queries by estimating the relative success of the user-oriented method.

7 Summary

Collaborative filtering through neighborhood-based interpolation ("kNN") is probably the most popular way to create a recommender system. The success of these methods depends on the choice of the interpolation weights, which are used to estimate unknown ratings from neighboring known ones. Nevertheless, the literature lacks a rigorous way to derive these weights. In this work we showed how the interpolation weights can be computed as a global solution to an optimization problem that precisely reflects their role. Comparison with past kNN methods on the Netflix data, demonstrated a significant improvement of prediction accuracy without a meaningful increase in running time.

Our second, complementary contribution is related to data normalization. Normalization is essential to kNN methods, as otherwise mixing ratings pertaining to different unnormalized users or items can produce inferior results. This work offered a comprehensive approach to data normalization, fitting 10 effects that can be readily observed in user-item rating data. Those effects cause substantial data variability and mask the fundamental relationships between ratings. Hence, their incluson brings the ratings closer and facilitates improved estimation accuracy.

A kNN method can be most effectively employed in an item-oriented manner, by analyzing relationships between items. However, certain cases call for a user-oriented approach. Examples include users that did not rate much or at all, or when it is needed to improve prediction accuracy by considering also user-user relationships. We show how to apply our method for the user-oriented case. Importantly, we offer a novel implementation, based on a low dimensional embedding of the users, that allows a significant computational speedup.

References

[1] G. Adomavicius and A. Tuzhilin, "Towards the Next Generation of Recommender Systems: A Survey of the State-of-the-Art and Possible Extensions", *IEEE Transactions on Knowledge and Data Engineering* **17** (2005), 634–749.

[2] R. M. Bell, Y. Koren and C. Volinsky, "Modeling Relationships at Multiple Scales to Improve Accuracy of Large Recommender Systems", *Proc. 13th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2007.

[3] J. Bennet and S. Lanning, "The Netflix Prize", KDD Cup and Workshop, 2007.

[4] B. Efron and C. Morris, "Data analysis using Stein's estimator and its generalization", *Journal American Statistical Association* **70** (1975), 311–319.

[5] S. Funk, "Netflix Update: Try This At Home", sifter.org/~simon/journal/20061211. html, 2006.

[6] D. Goldberg, D. Nichols, B. M. Oki and D. Terry, "Using Collaborative Filtering to Weave an Information Tapestry", *Communications of the ACM* **35** (1992), 61–70.

[7] K. Goldberg, T. Roeder, D. Gupta and C. Perkins, "Eigentaste: A Constant Time Collaborative Filtering Algorithm", *Information Retrieval* **4** (2001), 133–151.

[8] J. L. Herlocker, J. A. Konstan, A. Borchers and John Riedl, "An Algorithmic Framework for Performing Collaborative Filtering", *Proc. 22nd ACM SIGIR Conference on Information Retrieval*, pp. 230–237, 1999.

[9] J. Konstan, B. Miller, D. Maltz, J. Herlocker, L. Gordon and J. Riedl, "GroupLens: Applying Collaborative Filtering to Usenet News", *Communications of the ACM* **40** (1997), 77–87, www.grouplens.org.

[10] C.L. Lawson and B. J. Hanson, *Solving Least Squares Problems*, Prentice-Hall, 1974.

[11] G. Linden, B. Smith and J. York, "Amazon.com Recommendations: Item-to-item Collaborative Filtering", *IEEE Internet Computing* **7** (2003), 76–80.

[12] J. Nocedal and S. Wright, *Numerical Optimization*, Springer, 1999.

[13] R. Salakhutdinov, A. Mnih, and G. Hinton, "Restricted Boltzmann Machines for Collaborative Filtering", *Proc. 24th Annual International Conference on Machine Learning*, 2007.

[14] B. M. Sarwar, G. Karypis, J. A. Konstan, and J. Riedl, "Application of Dimensionality Reduction in Recommender System – A Case Study", *WEBKDD'2000*.

[15] B. Sarwar, G. Karypis, J. Konstan and J. Riedl, "Item-based Collaborative Filtering Recommendation Algorithms", *Proc. 10th International Conference on the World Wide Web*, pp. 285-295, 2001.